

Introduction to R

Valentina Staneva
vms16@uw.edu

June 29, 2015

R - The R Project for Statistical Computing

Main website: <http://www.r-project.org/>

- programming language based on S
- scripting language
- good for statistics and data analysis (but not limited to)
- free and open source

CRAN - The Comprehensive R Archive Network

Packages: <http://cran.us.r-project.org/>

Application Areas

Task Views: <http://cran.r-project.org/web/views/>

Some Links

- R intro
<http://cran.r-project.org/doc/manuals/R-intro.html>
- R Cookbook
<http://www.cookbook-r.com/>
- Quick-R
<http://www.statmethods.net/>
- R Cheat Sheets
<http://cran.r-project.org/doc/contrib/Short-refcard.pdf>
<http://cran.r-project.org/doc/contrib/Baggott-refcard-v2.pdf>

RStudio is an IDE environment for R.

Main components:

- R Console
- File editor
- Files, Plots, Packages, Help
- Workspace
- History

Installation

- 1 Installing R: <http://cran.us.r-project.org/>
- *try to remember the directory where you install R!*
- 2 Installing RStudio:
<http://www.rstudio.com/products/rstudio/download/>
- 3 Start RStudio
- *if needed, provide the R's installation path to RStudio*

Looking for help:

```
>help.start()
```

Setting your working directory:

```
>getwd()
```

```
>setwd("path_to_your_folder")
```

Basic types:

- numeric: integer - 1,3,-5; double - 0.5,-2.,3.5
- complex - i,1-2i
- character - "a", "ABBA", "?", "456"
- logical - TRUE/FALSE

Compound types:

- vectors (1D data)
- matrices (2D data)
- arrays (multi-dimensional data)
- lists (sequence of elements)
- dataframes (table-like data)
- factors (categorical data)

Creating Variables

Variable names can consist of letters, numbers, dot, underscore and need to start with a letter or dot not followed by a number.

```
a <- 1
b <- "a"
.number <- "4"
under_score <- "_"
```

- creating vectors

```
myNumericVector <- c(1,3,4)
myCharacterVector <- c("a","b","c")
myMixedVector <- c("a",1,"c")
```

Note: you can check class and type by:

```
class(myNumericVector)
typeof(myNumericVector)
```

Vectors need to contain elements of the same type!

Creating variables

- creating matrices

```
M<-matrix(c(1,2,3,4),nrow = 2,ncol = 2)
print(M)
```

```
##      [,1] [,2]
## [1,]    1    3
## [2,]    2    4
```

```
M<-matrix(c("h","w","e","o","l","r","l","l","o","d"),nrow = 2)
print(M)
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,] "h"  "e"  "l"  "l"  "o"
## [2,] "w"  "o"  "r"  "l"  "d"
```

Arrangement is columnwise!

More Ways to Create Vectors and Matrices

- sequence of integers:

```
v <- -1:5
```

- sequence of numbers:

```
v <- seq(from = 1,to = 10,by = 0.1)
```

- sequence of 100 zeros:

```
ones <- rep(0,100)
```

- sequence of random values:

```
r <- rnorm(10)
```

- 2x3 matrix of zeros:

```
M <- matrix(0,2,3)
```

- matrix with ones on the diagonal:

```
I <- diag(1,3)
```

- diagonal matrix:

```
D <- diag(c(1,2,3))
```

Concatenating Vectors and Matrices

- replicate a vector

```
vv <- rep(1:4,2)
```

- merge rows

```
M <- rbind(1:4,5:8)
```

- merge columns

```
M <- cbind(1:4,5:8)
```

- merge matrices

```
M<-cbind(M,M)
```

Elementwise Operations

$+$, $-$, $*$, $/$, $^$: addition, subtraction, multiplication, division, power
element by element for vectors or matrices of the same size.

Exercise:

- create two vectors of length 3
- what happens when you add one of those vectors to a 3x3 matrix?
- what happens when you add a vector of length 2 to the matrix?
- what happens when you add a vector of length 3 to 2x3 matrix?

Useful Functions

- `length(v)` - length of a vector
- `dim(M)` - dimensions of a matrix
- `length(M)` - length of a matrix
- `sum(v), sum(M)` - sum of all elements
- `min(v), max(M)` - min, max values
- `mean(v), mean(M)` - mean value
- `colSums(), rowSums(), colMeans(), rowMeans()`
- `var(), std()`
- `sqrt(), log(), exp(), sin(), cos(), ...`

Note: to learn how to use a function check out the documentation:

```
>help(function_name)
```

```
>?function_name
```

Accessing Values in Vectors and Matrices

- extracting individual elements: `v[3]`, `M[1,2]`, `M[5]`
- extracting a subvector: `v[1:3]`
- rearranging the elements: `v[c(3,1,2)]`
- extracting the odd elements: `v[seq(1,length(v),2)]`
- extracting a submatrix: `M[1:2,1:3]`
- extracting a row: `M[2,]`
- extracting a column: `M[,3]`
- extracting first and last column: `M[,c(1,dim(M)[2])]`

Logical Operations

We can perform logical operations on the whole vector/matrix

```
A <- matrix(rnorm(9),3,3)
```

```
A<0
```

```
##      [,1] [,2] [,3]  
## [1,]  TRUE FALSE  TRUE  
## [2,] FALSE FALSE  TRUE  
## [3,]  TRUE  TRUE FALSE
```

- negate statement

```
!A<0
```

- set the negative values to zero:

```
A[A<0] = 0
```

- find the locations of the zeros

```
which(A==0)
```

Data Frames

Data Frames are powerful data structures for data in a table form:

	mpg	cyl	dis	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1

Data Frame Example

R has some built-in datasets:

```
data()
data(mtcars)
summary(mtcars)
```

```
##           mpg           cyl          disp           hp
##  Min.      :10.40    Min.      :4.000    Min.      : 71.1    Min.      : 52.0
## 1st Qu.:15.43    1st Qu.:4.000    1st Qu.:120.8    1st Qu.: 96.5
## Median :19.20    Median :6.000    Median :196.3    Median :123.0
## Mean   :20.09    Mean   :6.188    Mean   :230.7    Mean   :146.7
## 3rd Qu.:22.80    3rd Qu.:8.000    3rd Qu.:326.0    3rd Qu.:180.0
## Max.   :33.90    Max.   :8.000    Max.   :472.0    Max.   :335.0
##           drat           wt           qsec           vs
##  Min.      :2.760    Min.      :1.513    Min.      :14.50    Min.      :0.0000
## 1st Qu.:3.080    1st Qu.:2.581    1st Qu.:16.89    1st Qu.:0.0000
## Median :3.695    Median :3.325    Median :17.71    Median :0.0000
## Mean   :3.597    Mean   :3.217    Mean   :17.85    Mean   :0.4375
## 3rd Qu.:3.920    3rd Qu.:3.610    3rd Qu.:18.90    3rd Qu.:1.0000
## Max.   :4.930    Max.   :5.424    Max.   :22.90    Max.   :1.0000
##           am           gear           carb
##  Min.      :0.0000    Min.      :3.000    Min.      :1.000
```


Useful Commands

- `colnames(mtcars)` - a vector of column names
- `rownames(mtcars)` - a vector of row names
- `dim(mtcars), nrow(mtcars), ncol(mtcars)` - dimensions of the data frame
- `length(mtcars)` - number of columns in the data frame
- `mtcars[5,4], mtcars[5:6,4:7]` - accessing elements
- `mtcars$mpg` - accessing columns
- `mtcars["Mazda RX4", "mpg"]`
- `mtcars[c("Mazda RX4", "Mazda RX4 Wag"),]`
- `mtcars[, c("hp", "mpg")]`

Note: rownames cannot be repeated, but colnames can ...

Creating Data Frames

We can have different types of variables in the columns.

- creating a data frame from vectors

```
> numbers = c(1,2,3)
> characters = c("a","b","c")
> df = data.frame(numbers,characters)
```

- creating a data frame from a matrix

```
> df = data.frame(M)
```

- creating an empty dataframe

```
> df =
data.frame(numbers=numeric(10),characters=character(10)))
```

- merging two data frames

```
df = rbind(df1,df2) - need same colnames and types!
```

```
df = cbind(df1,df2) - need same rownames!
```

- adding an extra column

```
df$new_name = values - need correct size
```

- removing a column

```
df$new_name <- NULL
```

Importing and Exporting Data

- store variables in .RData format (readable only by R)
- read, write spreadsheet-like files .csv (text file - very clean!)

```
data<-read.csv("matrix.csv")
```

```
write.csv(data,"matrix1.csv")
```

Importing and Exporting Data

- store variables in .RData format (readable only by R)
- read, write spreadsheet-like files .csv (text file - very clean!)

```
data<-read.csv("matrix.csv")
```

```
write.csv(data,"matrix1.csv")
```

Converting first column into row names:

```
data<-read.csv("matrix.csv",row.names = 1)
```

Using row names as an index:

```
write.csv(data,"matrix1.csv",row.names = FALSE)
```

FOR Loops

Example 1:

```
for (i in 1:10){  
  print(i)  
}
```

Example 2:

```
for (i in seq(1.5,10,0.5)){  
  print(i)  
}
```

Example 3:

```
for (i in c("a","b","c")){  
  print(i)  
}
```

IF Statements

Example 1:

```
for (i in 1:10){  
  if (i>5){  
    print(i)  
  }  
}
```

Example 2:

```
for (i in 1:10){  
  if (i>5 && i<8){  
    print("5<i<8")  
  }  
  else if (i<=5){  
    print("i<=5")  
  }  
  else{  
    print("i>=8")  
  }  
}
```

Functions

Functions in R can be stored as 'variables'.

```
myFunction <- function(a,b){  
  c = a + b  
  return(c)  
}
```

We can set default values:

```
myFunction <- function(a = 0,b = 0){  
  c = a + b  
  return(c)  
}
```

and then we don't have to specify all the values when in the call:

```
myFunction(3)  
  
## [1] 3  
  
myFunction(b=2)  
  
## [1] 2
```

Functions with multiple outputs

Rule: one can return only one object!

Case 1: outputs are of same type:

```
myFunction <- function(a,b){  
  c = a + b  
  d = a*b  
  result <-c(c,d)  
  names(result) = c("c","d")  
  return(result)  
}  
result <- myFunction(1,3)  
print(sprintf('The sum is %d and the product is %d.',  
              result["c"],result["d"]))  
  
## [1] "The sum is 4 and the product is 3."
```


Functions with multiple outputs

Rule: one can return only one object!

Case 2: outputs are of different type:

```
myFunction <- function(n){  
  v = rnorm(n)  
  M = matrix(rnorm(n*n),nrow = n)  
  result <-list(v = v,M = M)  
  return(result)  
}  
result <- myFunction(3)  
result$v  
result$M
```

A list is a sequence of any type of objects!

Apply Functions over Array Margins

`apply(X, MARGIN, FUN)` - apply function FUN to the first or second dimension of X

rowMeans:

```
apply(M, 1, mean)
```

colMeans:

```
apply(M, 2, mean)
```

You can easily vectorize your own functions!

`vapply` - for vectors

`lapply` - for lists

`sapply` - for vectors or lists

`mapply` - for multidimensional objects

`do.call` - similar to `lapply`

Exercise


1) create a function which takes a vector which might have missing values, and returns a vector where the missing values are substituted with the mean of the remaining values.

```
NA2Mean <- function(v){  
  #fill here  
  return(v_new)  
}  
NA2Mean(c(1,NA,3))  
c(1,2,3)
```

2) Apply your function to each column of data

Basic Plotting

```
x = seq(-1,1,0.1)
y = x^2
plot(x,y)
```



- plotting lines:
`plot(x,y,type = "l")`
- plotting points and lines:
`plot(x,y,type = "o")`

More options

- changing color

```
plot(x,y,col = "green")
```

- typical colors: "blue", "black", "red", "yellow", "orange", ...
`colors()`

- changing the symbol

```
plot(x,y,pch = 20)
```

- changing the linewidth

```
plot(x,y,lwd = 4)
```

More parameters

- `xlim = c(0,1), ylim = c(0,1)` (domain)
- `main, xlab, ylab` (the labels can be set with `title()`)
- `cex, cex.lab` (symbol size, label size)
- `text(), legend()`

Get parameters:



```
> par()
```

Set parameters:

```
> par(lwd = 5)
```

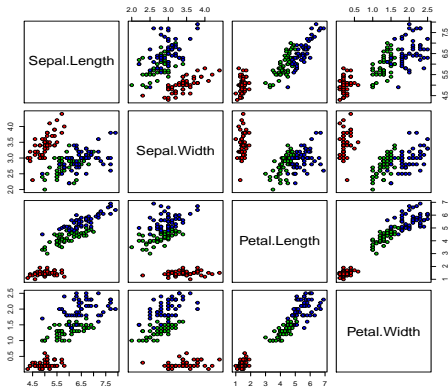
Iris Data Example

```
data(iris)
par(mfrow = c(2,1))
plot(iris$Sepal.Length,iris$Sepal.Width)
plot(iris$Petal.Length,iris$Petal.Width)
```



Iris Data Example

```
pairs(iris[1:4], pch = 21,  
bg = c("red", "green3", "blue")[unclass(iris$Species)])
```



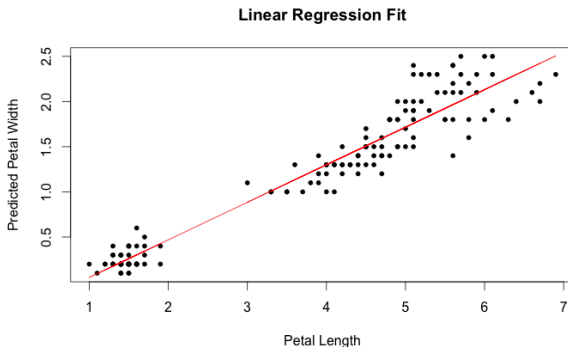
Iris Data Example

```
fit <- lm(Petal.Width ~ Petal.Length, data = iris)
summary(fit)

##
## Call:
## lm(formula = Petal.Width ~ Petal.Length, data = iris)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.56515 -0.12358 -0.01898  0.13288  0.64272
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.363076   0.039762  -9.131  4.7e-16 ***
## Petal.Length  0.415755   0.009582  43.387  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2065 on 148 degrees of freedom
## Multiple R-squared:  0.9271, Adjusted R-squared:  0.9266
## F-statistic: 1882 on 1 and 148 DF, p-value: < 2.2e-16
```



```
predicted<-predict(fit)
plot(iris$Petal.Length,iris$Petal.Width,pch = 20,
     xlab = "Petal Length",ylab = "Predicted Petal Width",
     main = "Linear Regression Fit")
lines(Petal.Length,predicted,col = "red")
```



Exporting graphs

```
png("myPlot.png")  
plot(x,y)  
dev.off()
```

Supports `png, jpeg, svg, pdf, postscript, ...`
You can set height, width, ...

Fancier graphs

- package: `ggplot2` - the Grammar of Graphics
- book: *R Graphics Cookbook*, by Winston Chung (2012)
(ebook in the library)
- tutorial: <http://www.cookbook-r.com/Graphs/>

Interactive graphs

- Plotly: <https://plot.ly/r/>
- rCharts: <http://ramnathv.github.io/rCharts/>
- Shiny apps: <http://shiny.rstudio.com>